

Searchable Symmetric Encryption: Improved Definitions and Efficient Constructions

Reza Curtmola, Juan Garay, Seny Kamara, Rafail Ostrovsky

Work appeared in CCS'06

Presented by: Rachit Garg and Aravind Birudu



Motivation

- Searching is the primary way we access data. We google stuff for most things today.
- We are outsourcing more and more of our data to third parties and we trust them less and less.
- Examples:

Industry: <https://numer.ai/>, for open source(data) machine learning

Governments: Aadhaar Database

Research Scientists: Genomic Data

The components

- Ways to encrypt data
- Model of the paper
- Prior Work
- Revisiting previous security definitions for SSE
- Two new notions of security for SSE
 - “Non-adaptive” security
 - “Adaptive” security
- New constructions
- Further work

Ways to compute on encrypted data

We know of six different ways to search on encrypted data, each based on one of the following cryptographic primitives:

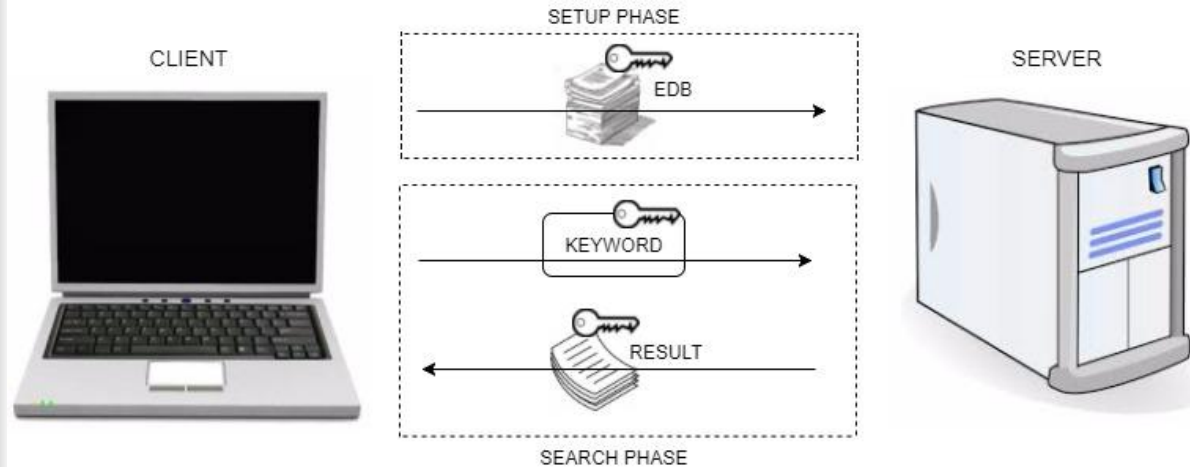
- property-preserving encryption
- functional encryption
- fully-homomorphic encryption
- searchable symmetric encryption
- oblivious RAMs
- secure two-party computation



The Model

Client uploads documents and an additional data structure to the server. This additional structure helps in searching in the database.

Goal: The leakage to the server should be minimum while the scheme being efficient



Security Definitions

- **History:** History of the queries, includes the document collection and the word queried.
- **Access Pattern:** Document indices returned by the queries.
- **Search Pattern:** Indicates what are the keywords that are searched for.
- **Trace:** Contains the length of the database, the access pattern and the search pattern.

ORAMs

- SSE through oblivious RAMs.
 - ORAM's can simulate data structures in a secure way, can support conjunctive queries.
 - Hides everything, even the access patterns.
 - Efficiency is logarithmic number of rounds for each read/write Lower bound shown by Goldreich and Ostrovsky.
 - Boneh-Kushilevitz-Ostrovsky-Skeith showed \sqrt{DB} communication with constant rounds.

Establishing a balance

Security



Efficiency



Trade-Offs

Get more efficiency! By leaking the access pattern but nothing else.

Prior Work

- Previous Attempts
 - “Practical techniques for searches on encrypted data” [SWP00]
 - “Secure Indexes” [Goh03]
 - “Privacy-preserving keyword searches on remote encrypted data” [CM05]
- [SWP00,Goh03,CM05]: “A secure SSE scheme should not leak anything beyond the outcome of a search”

[SWP00] : “any function of the plaintext that can be computed from the ciphertext can be computed from the length of the plaintext”

Issue: adversary gets to see search outcomes and search pattern

New Definition: “any function about the **documents** and the **keywords** that can be computed from the **encrypted documents**, the **index** and the **trapdoors** can be computed from the **length of the documents**, the **search outcomes** and the **search pattern**” (**adaptive and non adaptive!**)

IND2-CKA: indistinguishability against chosen-keyword attacks(used by [Goh03]) : “any function of the **documents** that can be computed from the **encrypted documents** and the **index** can be computed from the **length of the documents** and the **search outcomes**.

[CM05]: Any function that can be computed about the **documents** and **keywords** that can be computed from the **encrypted documents**, the **index** and the **trapdoors** can be computed from the **length of the documents** and the **search outcomes**. (**non adaptive**)

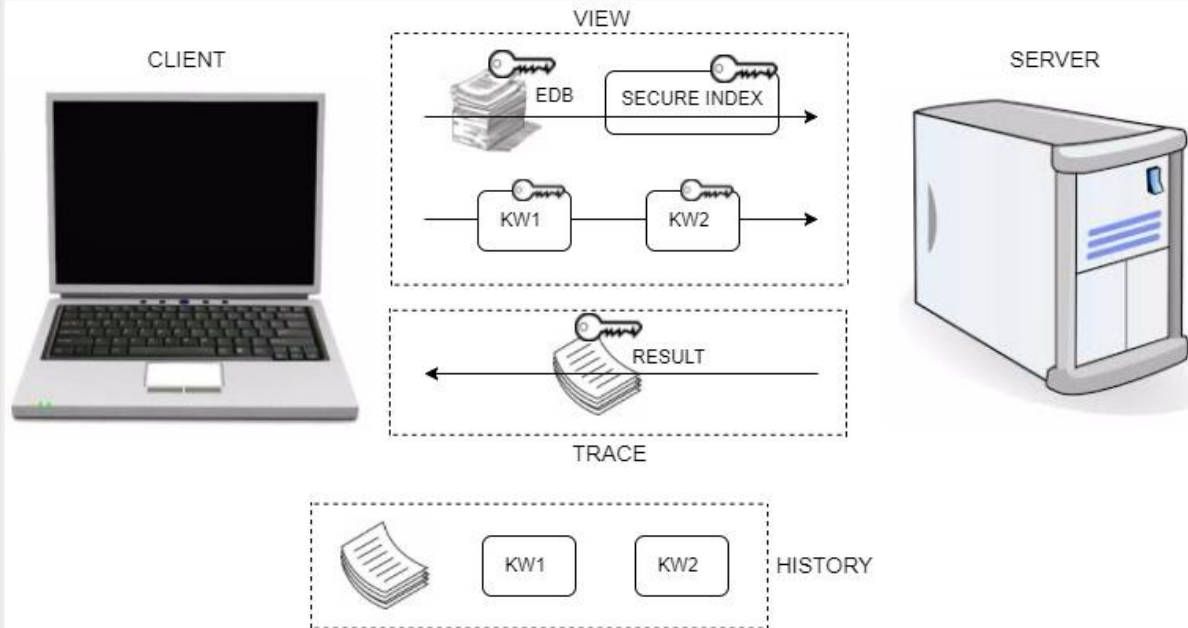
Question: Why not prove index secure in the sense of IND2-CKA and trapdoors “secure” using another definition?

They showed that there exists an SSE scheme that has IND2-CKA indexes and trapdoors that are “secure” but when taken together, adversary can recover keyword.

Algorithms

- **Keygen**(1^k): outputs symmetric key K
- **BuildIndex**($K, \{D_1, \dots, D_n\}$): outputs secure index I
- **Trapdoor**(K, w): outputs a trapdoor T_w
- **Search**(I, T_w): outputs identifiers of documents containing w (id_1, \dots, id_m)

SSE Model

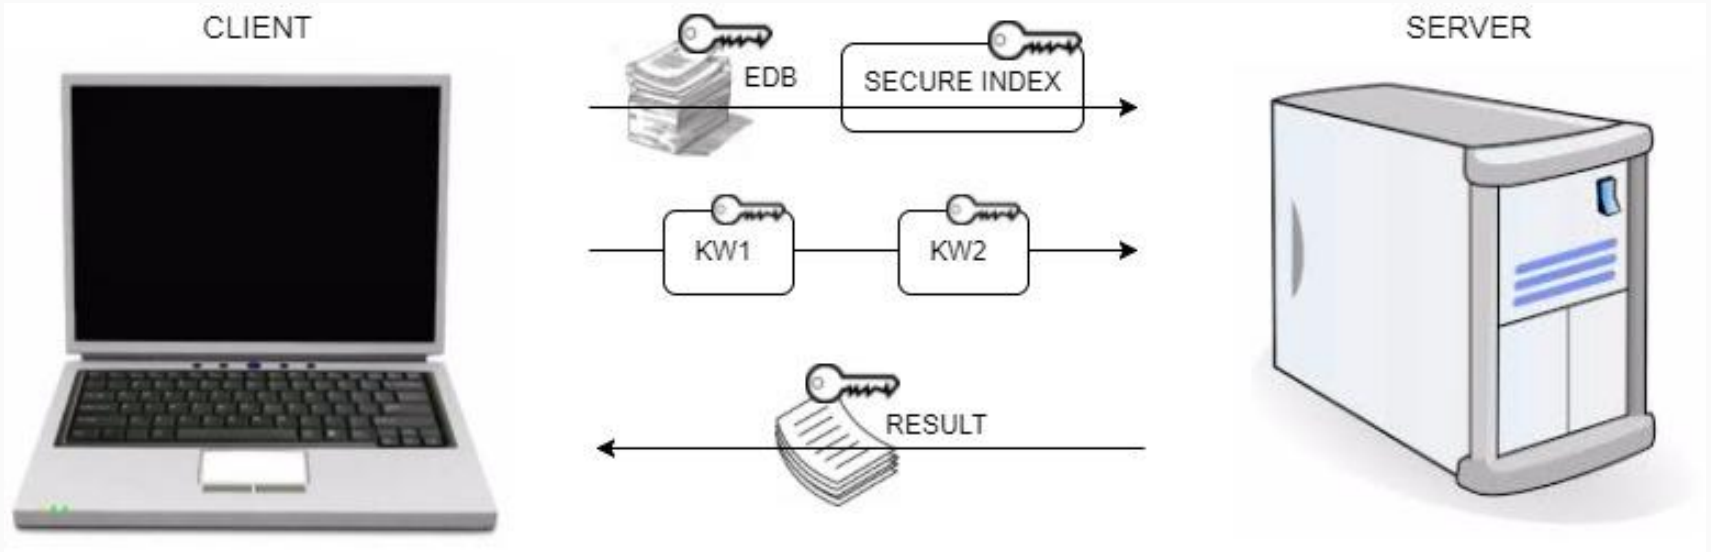


Adaptiveness

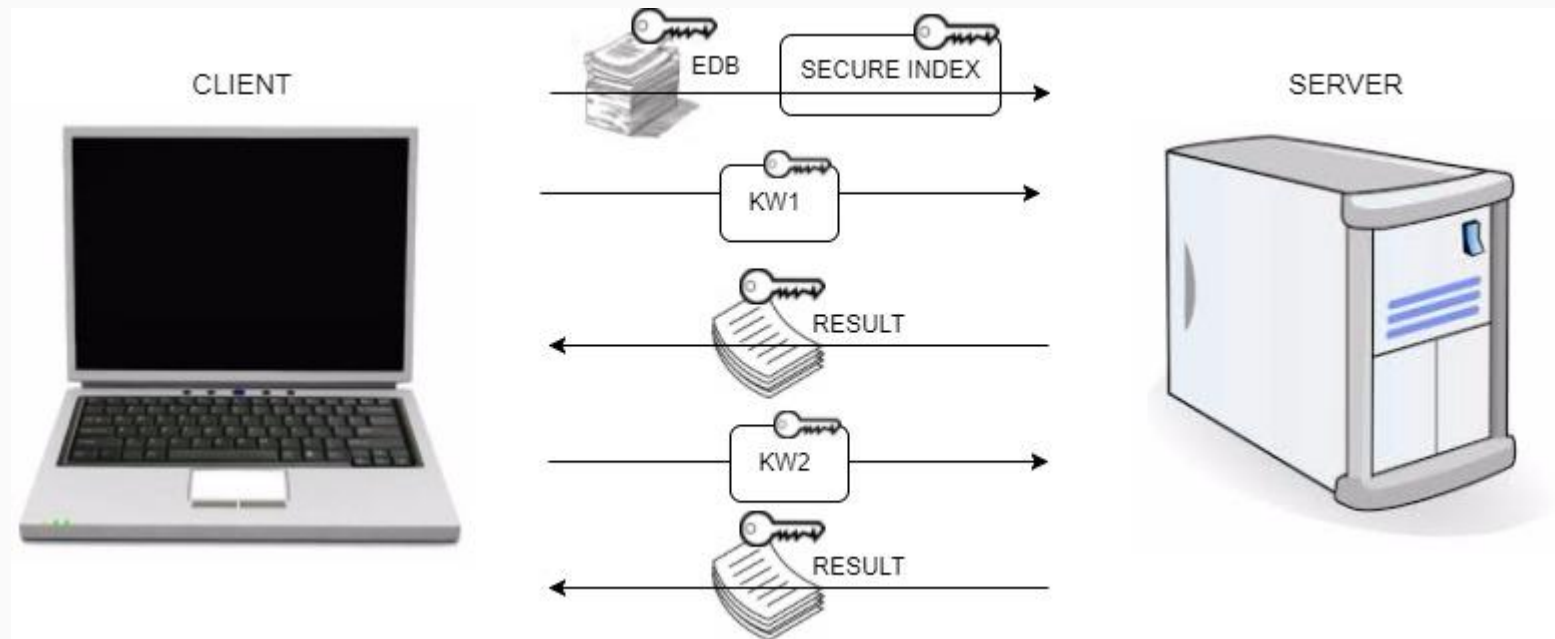
Non Adaptive: Non-adaptive adversaries make search queries without seeing the outcome of previous searches.

Adaptive: Adaptive adversaries can make search queries as a function of the outcome of previous searches.

NON ADAPTIVE



ADAPTIVE



Inverted Index Solution


- For each distinct keyword $w_i \in \delta(D)$, a linked list L_i is created that points to all documents containing that keyword.
- We then store all the nodes of all the lists in the array A permuted in a random order and encrypted with randomly generated keys.
- Before encrypting the j th node of list L_i , it is augmented with a pointer (with respect to A) to the $(j + 1)$ -th node of L_i , together with the key used to encrypt it.
- Note that by storing the nodes of all lists L_i in a random order, the length of each individual L_i is hidden.
- We then build a look-up table T that allows one to locate and decrypt the first node of each list L_i .

Inverted Index Solution

- The client generates both A and T based on the plaintext document collection D , and stores them on the server together with the encrypted documents.
- When the user wants to retrieve the documents that contain keyword w_i , it computes the decryption key and the address for the corresponding entry in T and sends them to the server.
- The server locates and decrypts the given entry of T , and gets a pointer to and the decryption key for the first node of L_i . Since each node of L_i contains a pointer to the next node, the server can locate and decrypt all the nodes of L_i , revealing the identifiers in $D(w_i)$.

Construction:- The Inverted Index Solution


► Building a Secure Index



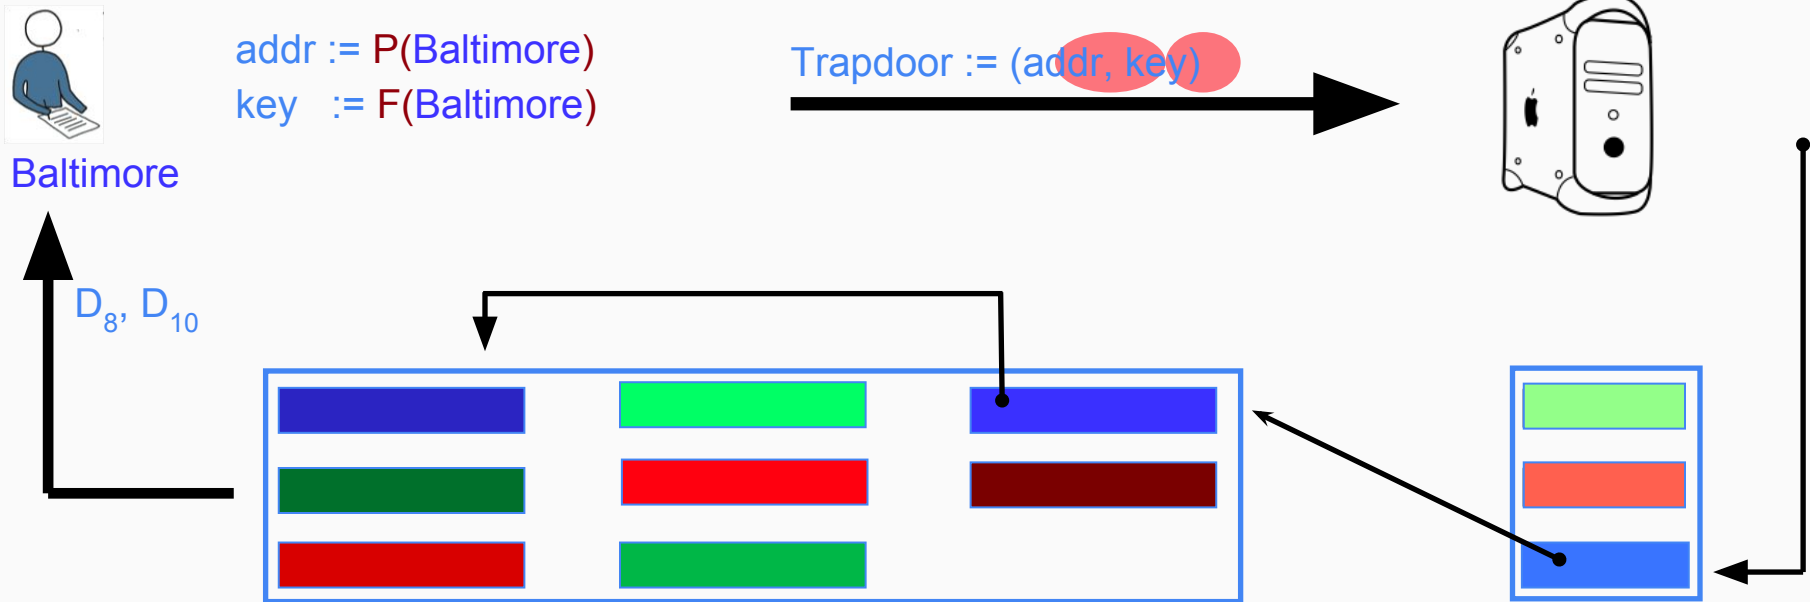
‣ Building a Secure Index

‣ P : PRP


‣ F : PRF



Searching



Multi - User SSE



Multi User SSE

- Indexes and trapdoors require same security notions as single-user SSE
- Revocation: owner can revoke searching privileges robust against user collusions
- Anonymity: server should not know who initiated search
- Simple construction that transforms single-user SSE schemes to multi-user SSE schemes.

Conclusion Table

Properties	ORAM	Other Solutions	This paper
Hides access patterns	yes	no	no
Server computation	$O(\log^2 n)$	$O(n)$	$O(1)$
Server storage	$O(n \log n)$	$O(n)$	$O(n)$
Number of rounds	$\log n$	1	1
communication	$O(\log^2 n)$	$O(1)$	$O(1)$
Adaptive adversaries	yes	no	no

Further Work

- This was a foundational work in SSE, it defined the necessary security definitions. There have been several improvements to the solution described in this paper.
- The use of FKS dictionaries have been removed.
- The inverted index solution is a static scheme, dynamic schemes have been proposed.
- Cash, Jarecki, Jutla, Krawczyk, Rosu and Steiner have extended inverted index solution to handle boolean queries.

Thanks!